Buxin Su

Curriculum Vitae

Philadelphia, Pennsylvania | subuxin@sas.upenn.edu | Personal Website | Google Scholar

EDUCATION

Ph.D. in Applied Mathematics and Computational Science 2023.09 – 2027.06 (Expected)
University of Pennsylvania
Advisors: *Bingxin Zhao* and *Weijie Su*

M.S. in Pure Mathematics 2021.09 – 2023.06
University of Pennsylvania
Advisors: *Bingxin Zhao* and *Robin Pemantle*
Thesis: *The exact risks of reference panel-based regularized estimators*

B.S. in Pure Mathematics & **Applied Mathematics Specialist** 2017.09 – 2021.06
University of Toronto
Graduated with High Distinction

SELECTED AWARDS AND HONORS

NESS Student Research Award 2026.05
New England Statistics Symposium

Two Sigma Ph.D. Fellowship Finalist 2026.01
Two Sigma

Lawrence D. Brown Best Paper Award 2025.10
The Wharton School, University of Pennsylvania

Discussion Paper, Journal of the American Statistical Association 2025.08
American Statistical Association

ICSA NSF Travel Award 2025.06
ICSA 2025 Applied Statistics Symposium

Student Paper Award, Statistical Learning and Data Science Section 2024.01
American Statistical Association

RESEARCH INTERESTS

My research lies at the intersection of statistics and machine learning (ML), with applications in the social sciences. I focus on the reliability and societal implications of artificial intelligence (AI). I develop statistical theory and methods to address three connected challenges: improving how the AI community evaluates research, increasing the utility of social data under privacy constraints, and making statistical learning more reliable for biomedical applications. Contributions include **methods incorporated into the ICML 2026 review process**, methods that could **reduce noise variance in the 2020 U.S. Census by about 20%** while maintaining comparable privacy protection, and studies of the statistical foundations of biomedical prediction.

Statistical Methods for Evaluating AI Research

- **Using Author Knowledge to Improve Peer Review at AI Conferences.** Developing, evaluating, and implementing statistical methods that incorporate authors' knowledge of their own work, specifically their rankings of their own papers, into peer review at AI/ML conferences [J1, J3, J4, J5]. ICML 2026 used authors' self-rankings to flag submissions with potentially low-quality reviews for further assessment by Area Chairs.
- **Randomized Experiment on AI-Assisted Scientific Review.** Evaluating how policies on AI assistance affect review behavior and outcomes through a randomized experiment and survey at ICML 2026, a conference involving over 17,000 reviewers [C3].

Improving the Utility of Social Data under Privacy Constraints

- **Exact Privacy Accounting for Social Data.** Characterizing the privacy–utility trade-off of the 2020 U.S. Census disclosure-avoidance system through theoretical analysis and showing that it added more noise than necessary: noise variances could be reduced by 15.08%–24.82% while maintaining comparable privacy protection [J2]. Developing practical methods to determine precisely the noise required for large-scale census releases [P3].
- **Private Learning with Improved Utility.** Developing theoretical tools and algorithms to improve learning utility while preserving privacy [C1, P5] and to improve predictive performance in federated learning under privacy constraints [C2].

Reliable Statistical Learning for Biomedical Applications

- **Characterizing Prediction Risk in Statistical Learning.** Characterizing the prediction risks of statistical learning methods commonly used in biomedical applications when trained on summary data in high-dimensional settings [P1, P2]. Complementary work examines how synthetic rationale data used in fine-tuning affect disease prediction [P4].

SELECTED PUBLICATIONS

** denotes equal contribution.*

Journal Articles (Published or Accepted)

J1. **Buxin Su**, Jiayao Zhang, Natalie Collina, Yuling Yan, Didong Li, Kyunghyun Cho, Jianqing Fan, Aaron Roth, and Weijie Su, "The ICML 2023 Ranking Experiment: Examining Author Self-Assessment in ML/AI Peer Review," *Journal of the American Statistical Association*, 2025. *Link* **[Selected as a Discussion Paper]**

Discussants: Dinghuai Zhang and Yoshua Bengio; Andreas Haupt and Sanmi Koyejo; Xiao-Li Meng; Weichen Wang and Chengchun Shi; Ying Wang and Mengye Ren; Linjun Zhang and Lexin Li.

J2. **Buxin Su**, Weijie Su, and Chendi Wang, "The 2020 US Decennial Census Is More Private Than You (Might) Think," *Proceedings of the National Academy of Sciences*, 2025. *Link* **[Won the 2025 Lawrence D. Brown Best Paper Award] [Featured by *New Scientist*]**

J3. **Buxin Su**, Natalie Collina, Garrett Wen, Didong Li, Kyunghyun Cho, Jianqing Fan, Bingxin Zhao, and Weijie Su, "Self-Rankings as a Predictor of Scientific Impact Beyond Peer Review," *Nature Computational Science*, 2026. *Link* **[Featured by *Nature Index*, *Nature Computational Science News & Views*, *Wharton AI & Analytics Initiative*, *Penn Today*, and *The Daily Pennsylvanian*]**

J4. **Buxin Su**, Jiayao Zhang, Natalie Collina, Yuling Yan, Didong Li, Kyunghyun Cho, Jianqing Fan, Aaron Roth, and Weijie Su, "Rejoinder: The ICML 2023 Ranking Experiment: Examining Author Self-Assessment in ML/AI Peer Review," *Journal of the American Statistical Association*, 2026. *Link*

J5. Garrett Wen, **Buxin Su**, Natalie Collina, Zhun Deng, and Weijie Su, "Recommending Best Paper Awards for ML/AI Conferences via the Isotonic Mechanism," *Journal of Machine Learning Research*, accepted, 2026. *Link*

Conference and Workshop Papers

C1. Chendi Wang*, **Buxin Su***, Jiayuan Ye, Reza Shokri, and Weijie Su, "Unified Enhancement of Privacy Bounds for Mixture Mechanisms via f-Differential Privacy," *Neural Information Processing Systems (NeurIPS)*, 2023. *Link*

C2. Xiang Li*, **Buxin Su***, Chendi Wang*, Qi Long, and Weijie Su, "Mitigating the Privacy–Utility Trade-off in Decentralized Federated Learning via f-Differential Privacy," *Neural Information Processing Systems (NeurIPS Spotlight)*, 2025. *Link*

C3. Sunnie S. Y. Kim*, Wesley Hanwen Deng*, Jennifer Wortman Vaughan, **Buxin Su**, Weijie Su, Alekh Agarwal, Sharon Li, Martin Jaggi, Daniel G. Goldstein, Nihar B. Shah, and Miroslav Dudík, "Use and Effects of LLMs in Peer Review: A Randomized Experiment and Survey at ICML 2026," *NeurIPS 2026 Workshop on AI-Native Academia*, 2026. *Link* [**Won the Best Paper Award**]

Preprints

P1. **Buxin Su**, Qiang Sun, Xiaochen Yang, and Bingxin Zhao, "The Exact Risks of Reference Panel-Based Regularized Estimators," under major revision at the *Journal of the American Statistical Association*, 2024. *Link* [**Won the 2024 American Statistical Association Student Paper Award, Statistical Learning and Data Science Section**]

P2. **Buxin Su**, Jiaoyang Huang, Jin Jin, and Bingxin Zhao, "On Self-Training of Summary Data with Genetic Applications," under major revision at the *Journal of the Royal Statistical Society Series B*, 2025. *Link* [**Won the 2026 New England Statistics Symposium Student Research Award**]

P3. **Buxin Su**, Weijie Su, and Chendi Wang, "A Sieve-Accelerated Quadrature Method for Exact Privacy Accounting in the 2020 U.S. Decennial Census," *arXiv:2606.29835*, 2026. *Link*

P4. **Buxin Su**, Bingxuan Li, Cheng Qian, Yiwei Wang, Jin Jin, and Bingxin Zhao, "Supervised Fine-tuning with Synthetic Rationale Data Hurts Real-World Disease Prediction," *arXiv:2606.10279*, 2026. *Link*

P5. **Buxin Su***, Qiaoshi Yang*, Yiding Su*, and Chendi Wang, "Unifying Privacy Accounting: Information Equivalence and Information Loss," *arXiv:2610.02414*, 2026. *Link*

PROFESSIONAL SERVICE

Associate Chair
Forty-Third International Conference on Machine Learning (ICML 2026) *2026.01 – 2026.07*

Session Chair
New England Statistics Symposium 2026 *2026.05*
Session: "Modern Statistical Learning, Predictive Inference, and Uncertainty Quantification"

Joint Statistical Meetings 2025 *2025.08*
Session: "Large Language Models in Biomedical and Statistical Knowledge Discovery"

Journal Reviewer
IEEE Access, Behaviour & Information Technology, IEEE Journal on Selected Areas in Information Theory

Conference Reviewer
International Conference on Learning Representations (ICLR 2026)

TEACHING

Teaching Assistant
STAT 1020 – Introductory Business Statistics *2026.09 – 2026.12*
The Wharton School, University of Pennsylvania

INVITED TALKS

Statistics for AI and Society: Evaluating AI Research and Improving Social Data 2026.10
Statistics Seminar, New Jersey Institute of Technology

On Self-Training of Summary Data with Genetic Applications 2026.05
New England Statistics Symposium 2026

The Three-Year ICML Ranking Experiment: Using Authors' Self-Rankings to Improve Peer Review 2026.01
Two Sigma Ph.D. Fellowship Finalist Reception 2026

The 2020 U.S. Decennial Census Is More Private Than You Think 2025.10
Lawrence Brown Student Workshop 2025

The 2020 U.S. Decennial Census Is More Private Than You Think 2025.06
ICSA Applied Statistics Symposium 2025

The Exact Risks of Reference Panel-Based Regularized Estimators 2024.08
Joint Statistical Meetings 2024

The ICML 2023 Experiment: Can Authors' Opinions on Their Own Submissions Assist Peer Review? 2023.11
Warren Center Lightning Talks, University of Pennsylvania

The Exact Risks of Reference Panel-Based Regularized Estimators 2023.07
ICSA 2023

LANGUAGES

- English: Fluent
- Mandarin Chinese and Cantonese: Native

REFEREES

Weijie Su
Professor
Department of Statistics and Data Science
The Wharton School
University of Pennsylvania

411 Academic Research Building
265 South 37th Street
Philadelphia, PA 19104
suw@wharton.upenn.edu

Bingxin Zhao
Associate Professor
Department of Statistics and Data Science
The Wharton School
University of Pennsylvania

413 Academic Research Building
265 South 37th Street
Philadelphia, PA 19104
bxzhao@wharton.upenn.edu

Jianqing Fan
Frederick L. Moore '18 Professor of Finance
Professor of Statistics and Machine Learning
Department of Operations Research
and Financial Engineering
Princeton University

205 Sherrerd Hall
Princeton, NJ 08544
jqfan@princeton.edu

Kyunghyun Cho
Glen de Vries Professor of Health Statistics
Professor of Computer Science and Data Science
Department of Computer Science
and Center for Data Science
New York University

60 Fifth Avenue, Room 616
New York, NY 10011
kyunghyun.cho@nyu.edu